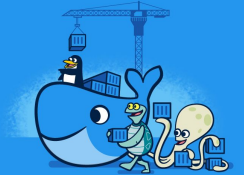

DevOps

Docker Containers

Containers

What is a Container ?

- Lightweight portable package with binaries allowing to encapsulate and run a program.

- Let’s say you want to run a program. All you need is :
    - A minimal base image allowing your to install / operate your program
    - A Runtime for the program
    - Dependencies
    - Your Application’s code

- You don’t need to bother with your Kernel and what’s installed on it.

- In a project, containers are very practical deliverable since they can be run “anywhere”

Containers

Containers vs VMs ?

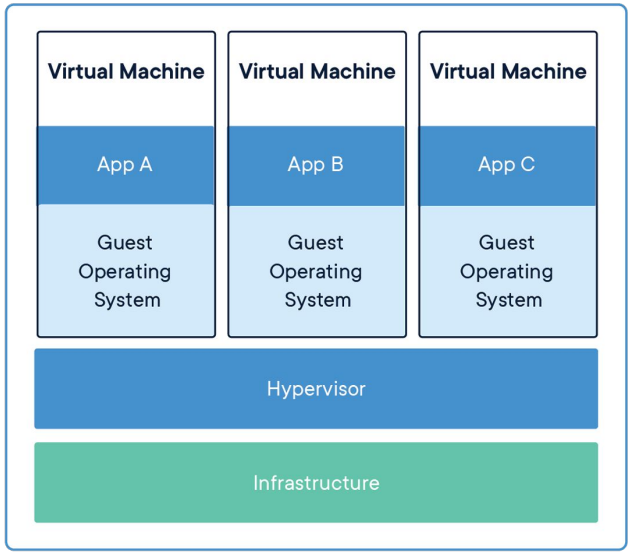

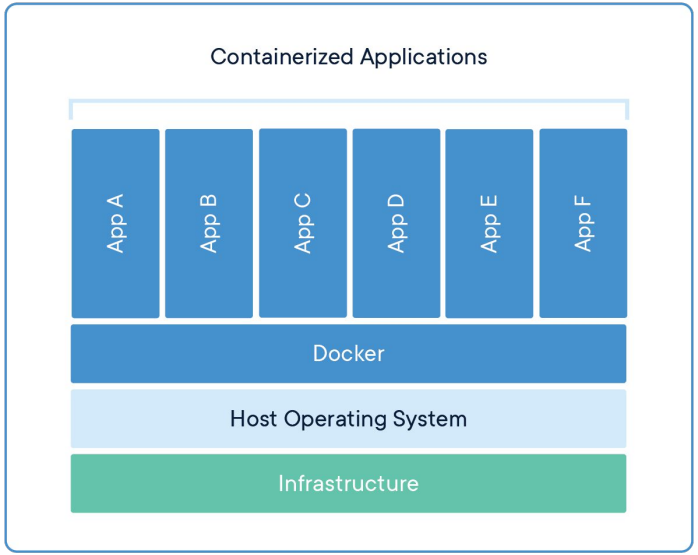

source : docker.com

- Hardware level
- Whole OS
- Reserves resources
- Slow to start

- Lightweight
- Share computer's OS & resources
- Start very fast

Containers

Why Containers / Advantages

- Start instantly
- Easy to deploy & run
- Easy to reproduce problems
- Easy to scale & distribute
- Built from a declarative state & reproductible
- Many available tools ready to use (public images)
- Platform independent

Containers

Orchestration

- Once we have containers, it is essential to:
  - Handle auto-scaling / replication.
  - Check their health & restart them.
  - Ensure they can communicate together and access persistent storage volumes.

- Container Orchestration tools come to the rescue:
  - Mesosphere Marathon
  - Kubernetes
  - Docker Swarm

Reminders on Containerization

- Containers allow deploying multiple services on hosts while isolating them from each other and from hosts packages.
- Consistent way to deliver and run applications
- Fast deployment, easy to scale
- Offer modularity to operate and maintain services

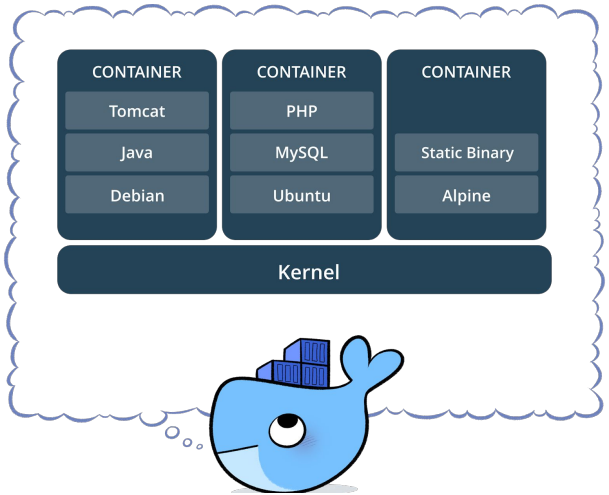

Reminders on Containerization

Best practices with Containers:

- Each container should do one thing and do it well.
- Containers should be as light as possible.
- Users should not have too much privileges (disable root user).
- Scan non-official images before running them.

Docker

- Created for dotCloud’s needs by Solomon Hykes
- Distributed as an Open Source project since 2013
- Written in Go
- Available on Linux, macOS and Windows since 2020 (Windows Subsystem for Linux 2)

Docker Components

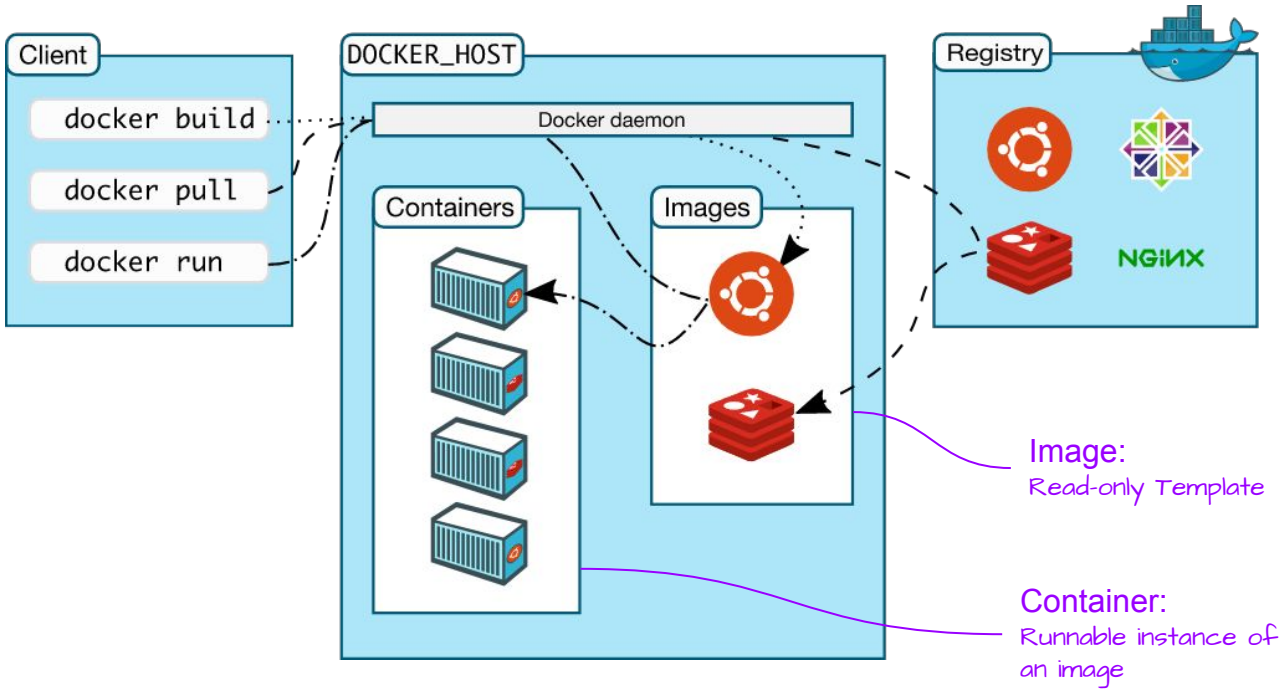

Docker Images

- Images are built with stacked layers (increments)

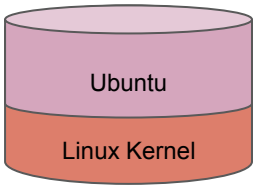

Ubuntu Image

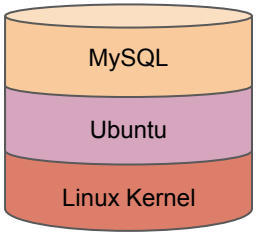

MySQL on Ubuntu Image

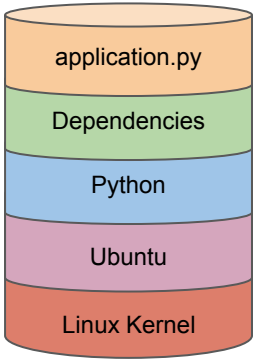

Some Python app on Ubuntu Image

Docker Images

- Images are fully-named and referenced with :
  - an optional prefix, the registry hostname
  - a name identifying the image
  - a tag (version), defaulting to “latest”

nginx refers to the latest tag for the official nginx image.

elasticsearch:7.14.1 refers to the official elasticsearch image with version 7.14.1

bitnami/kafka:2.8.1 refers to Kafka version 2.8.1 built by bitnami

Docker Images

- Images are built using a base image and performing actions on it.
- Either described in Dockerfiles or built by committing changes to an image.

Image name with tag

```
FROM ubuntu:20.04

RUN apt-get update && apt-get install -y python3 pip

RUN pip install elasticsearch

ADD application.py /opt/application.py

ENTRYPOINT python /opt/application.py
```

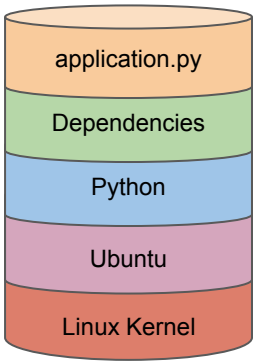

Docker Images

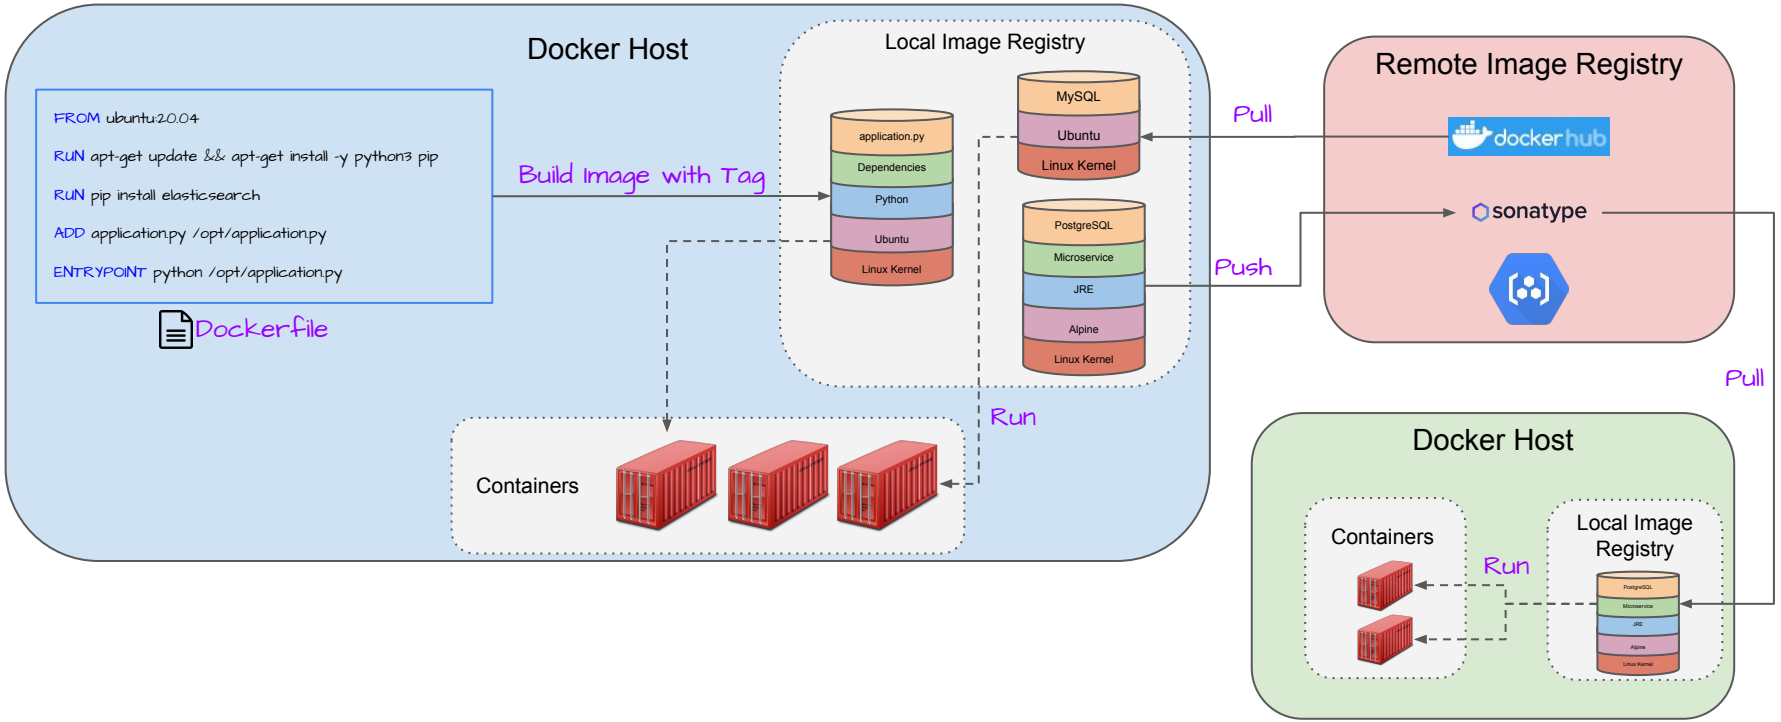

Docker Containers

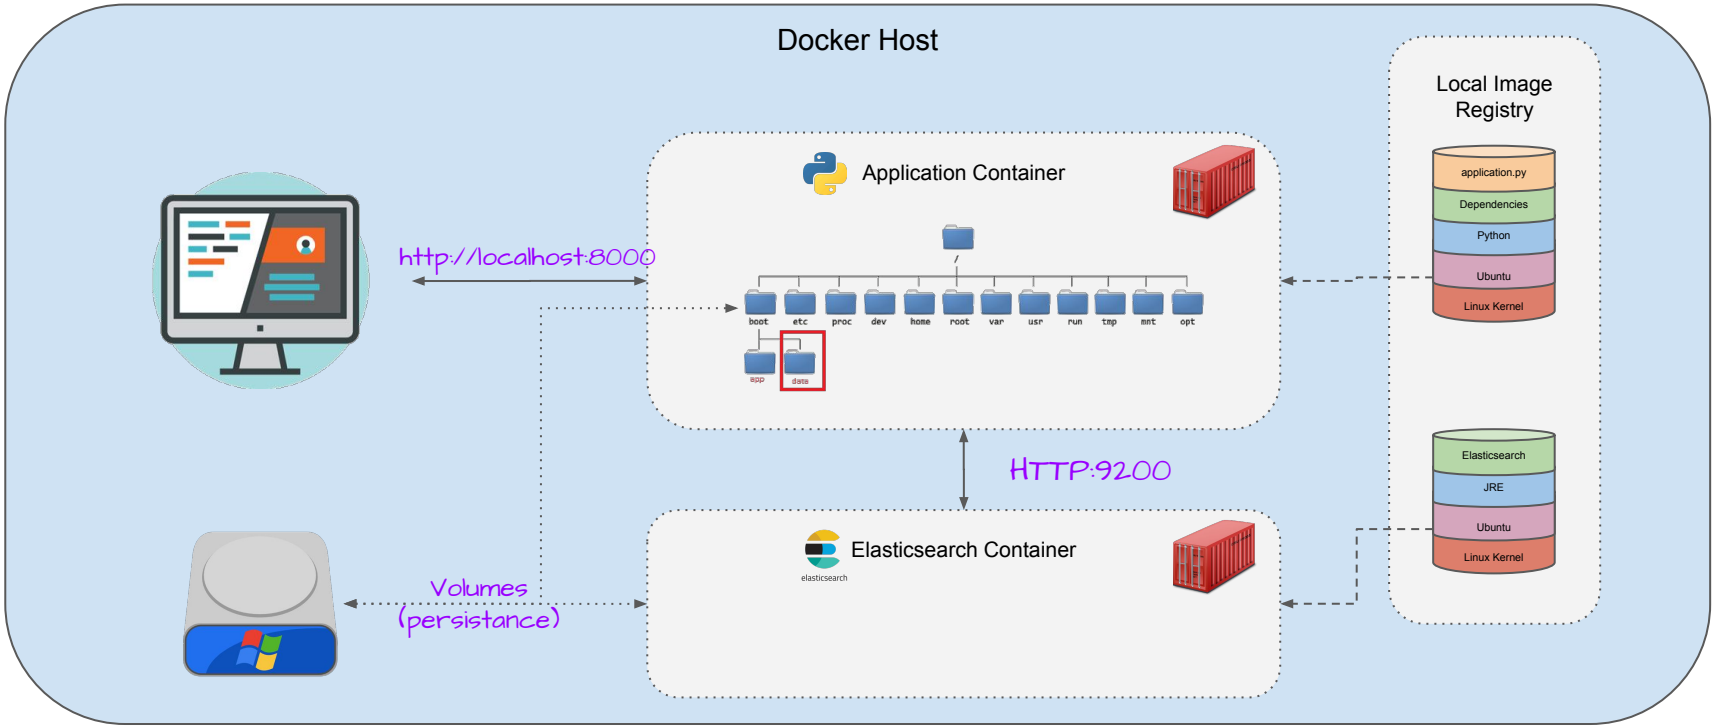

Docker Containers

Lifecycle

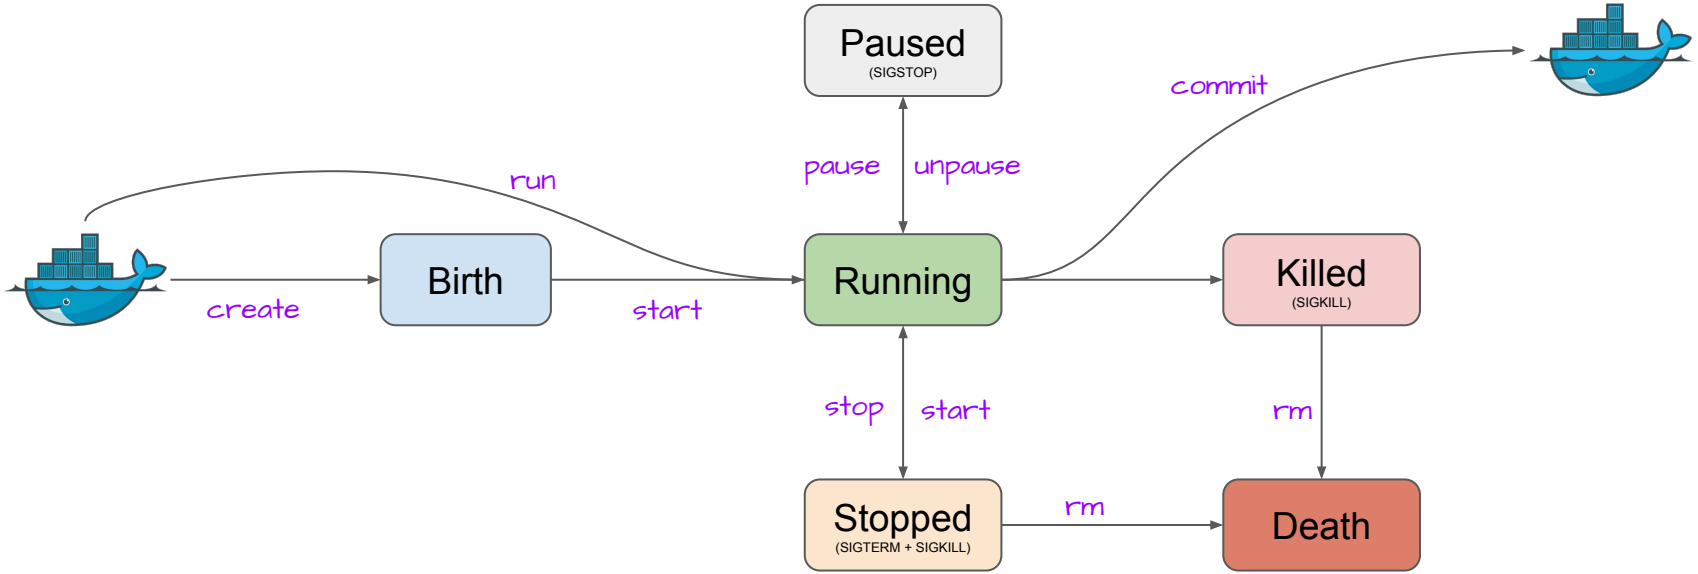

Docker Network

- Docker comes with the ability to create Virtual Networks:
  - Allow you to interact with Containers
  - Allow containers to communicate

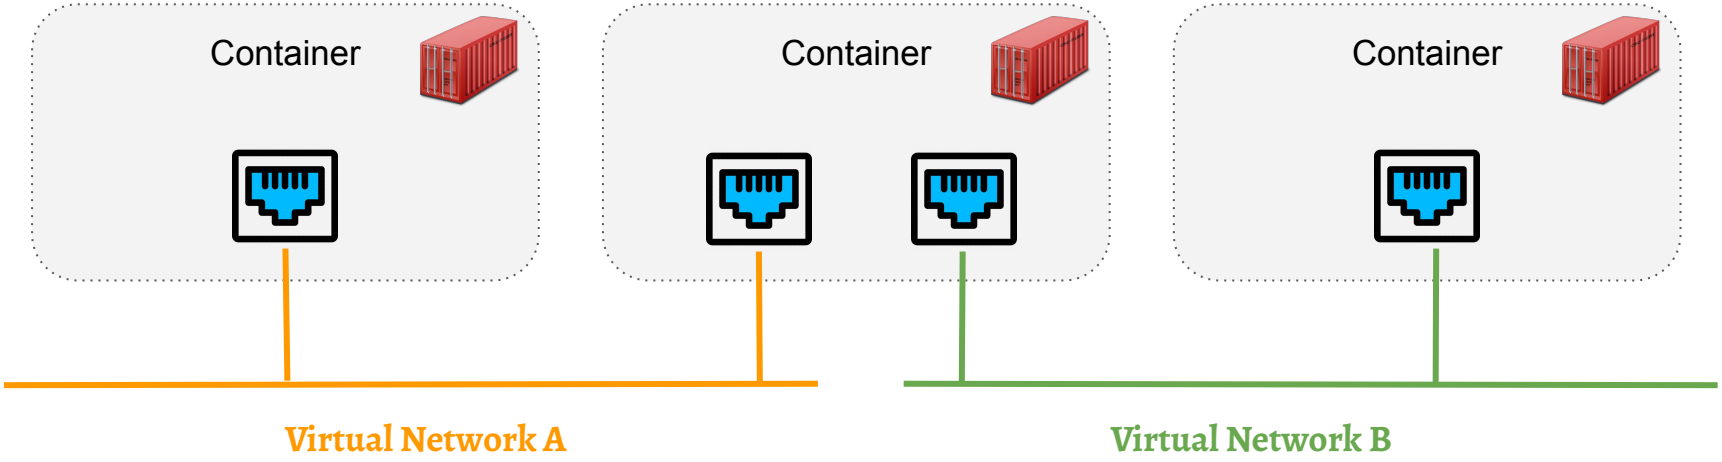

Docker Network

- Multiple network isolation possibilities : nothing / everything / virtual networks
- Allow exposing / mapping specific ports

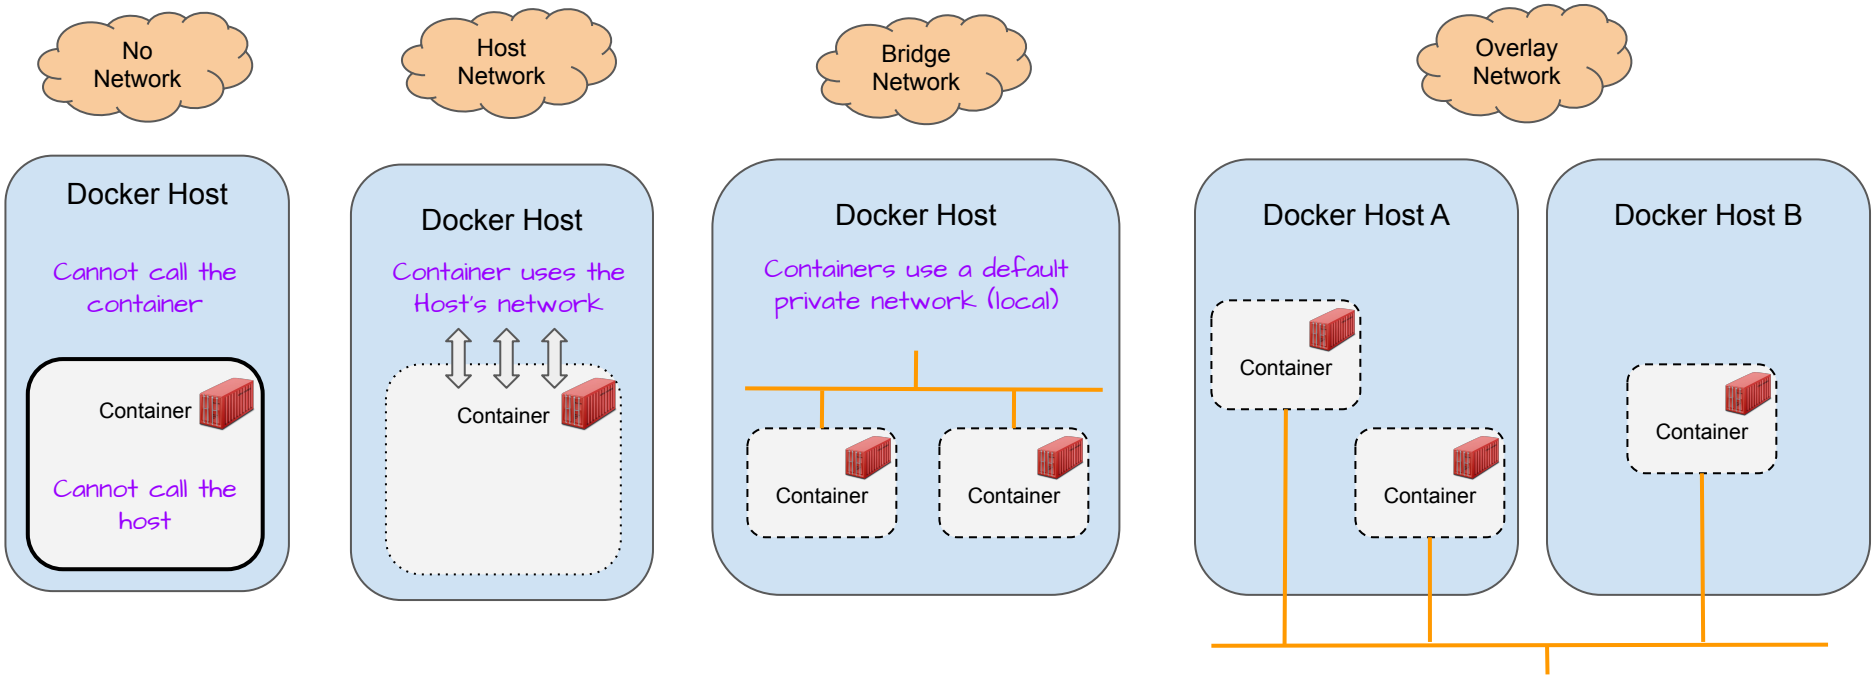

Docker Network

- Containers can call each other using their names using an internal DNS.

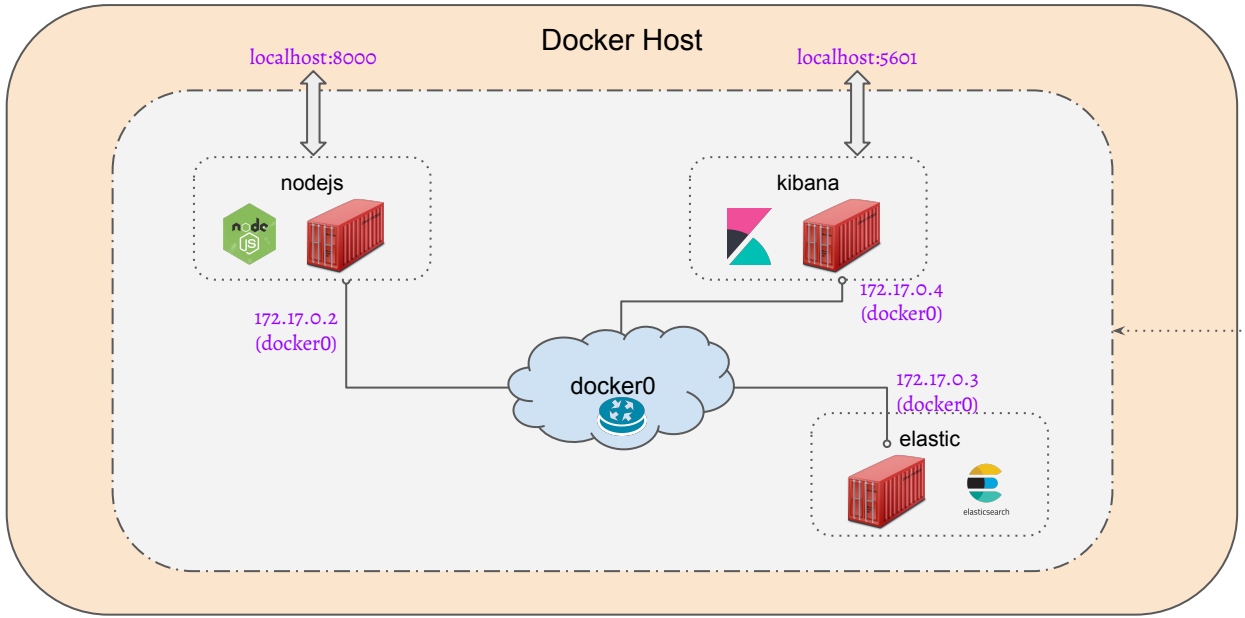


```
{
   "Name": "my-network",
   "Driver": "bridge",
   "IPAM": {
      "Driver": "default",
      "Config": [
         {
            "Subnet": "172.17.0.0/16",
            "Gateway": "172.17.0.1"
         }
      ]
   },
   "Containers": {

"35a42839b835391bb4ba021ffd1d151593a17663edd75f7d64563fb65c025037": {
         "Name": "nodejs",
         "MacAddress": "02:42:ac:12:00:02",
         "IPv4Address": "172.17.0.2/16"
      },

"c0a0f3cc90da4372ad2f3f82cf2a7cfda4f64c7f2513170f870792823b068d97": {
         "Name": "elastic",
         "MacAddress": "02:42:ac:12:00:03",
         "IPv4Address": "172.17.0.3/16"
      }
   }
}
```

Docker Network

Network Mode	Description	Example Use case
none	No networking. Cannot be called from outside and cannot call external services.	Standalone worker 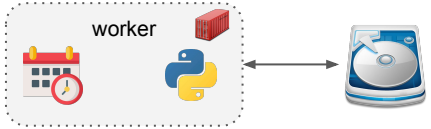
bridge	**Default**. Creates virtual interfaces to allow communication between the containers within the network. This virtual network is shared with Host.	We may need to : - Connect containers together - Expose specific ports - Access the external world 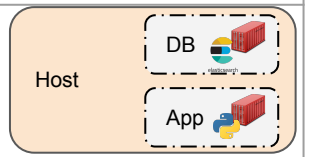
host	Containers use the Host's network. Everything's available from inside and outside. There is **no network-level isolation**.	All we're interested in is the environment isolation (binaries / packages & versions). Services are managed at the container level. 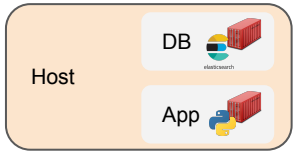
overlay	Create a Virtual network allowing containers to see each other across the L2 physical network.	Containers: - Are in a same subnet. - Can access external world - Can be distributed across multiple nodes (orchestration) 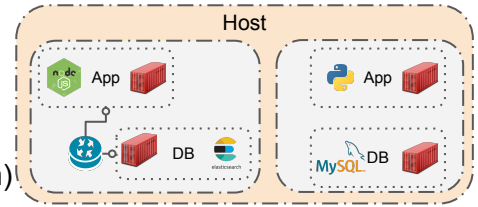

Docker

Testing with Containers

- Images exist for multiple services.
- Possible to create custom images.
- Containers are lightweight and can be run everywhere.
- Containers allow to **reproduce** things.
- Containers can be used for Integration tests.

https://www.testcontainers.org/

Docker Commands

Images

Command	Description	Examples
docker pull <image name>[:tag]	Downloads an image with a specific tag (latest by default).	docker pull ubuntu docker pull elasticsearch:7.14.1
docker build	Builds an image using a Dockerfile.	docker build -t myapp:0.1 .
docker images	Lists images with tags available in the local image registry.	docker images docker images --all
docker commit <containerId> <image>	Creates a new image from a container.	docker commit 597ce1600cf4 custom-ubuntu:20.04
docker rmi <image>[:tag]	Removes an image from the local registry.	docker rmi myapp:0.1
docker push <image name>[:tag]	Publishes an image from the local registry to a remote.	docker image push registry-host:port/myapp:0.1
docker save	Exports the image to a tar archive.	docker save myapp:0.1 > myapp_0-1.tar
docker load	Imports an image from a tar archive.	docker load < myapp_0-1.tar

Docker Commands

Containers

Command	Description	Examples
docker run <image name>[:tag]	Creates a container from an image.	docker run -it ubuntu:20.04 docker run -e "discovery.type=single-node" -d -p 9200:9200 elasticsearch:7.14.1 docker run -v /opt/app:/app --link elastic:elastic faucet/python3 pyhton /app/main.py
docker exec	Uses a container and runs actions on it.	docker exec -it 9e8eb83f233f ls /app
docker ps	Lists containers on host	docker ps -a
docker logs <containerId>	Shows the console logs for a specific container.	docker logs 9e8eb83f233f
docker stop <containerId>	Stops the specified container(s).	docker stop 9e8eb83f233f ad14edb7724a
docker rm <containerId>	Removes the specified container(s). The container(s) must be stopped or killed before they can be removed.	docker rm 9e8eb83f233f ad14edb7724a

Dockerfile Instructions

https://docs.docker.com/engine/reference/builder/

Instruction	Description	Examples
FROM <image>[:tag]	Used to specify a base image. Any Dockerfile must start with a FROM instruction, unless arguments are required.	**FROM** ubuntu:20.04
RUN <cmd> [args]	Executes any command / program available at this point. May also include arguments.	**FROM** ubuntu:20.04 **RUN** apt-get update && apt-get install python3 **RUN** python3 myscript.py
ARG	Allows to define arguments, and optionally affect a default value. Build-time variables. If no value set, the image won't be built.	**ARG** UBUNTU_VERSION=20.04 **FROM** ubuntu:$UBUNTU_VERSION
ENV	Defines environment variables. Available both at build and run time.	**FROM** ubuntu:20.04 **ADD** application /opt/application **ENV** HTTP_PORT=8000
ADD	Lets you copy files from your host (or an URL) to a specified destination path.	**FROM** ubuntu:20.04 **ADD** application /opt/application **ADD** https://someserver.com/archive.tar.gz /opt/
EXPOSE <port>	Informs Docker that the containers should listen at the specified ports. It does not actually "publish" the port at runtime.	**FROM** ubuntu:20.04 **ADD** application /opt/application **EXPOSE** 8000
ENTRYPOINT	Specifies what will be executed when creating a container.	**FROM** ubuntu:20.04 **ADD** application /opt/application **ENTRYPOINT** /opt/application/main.py 8000

Dockerfile

Advices

- Use official images whenever possible:
  - Clear documentation
  - Maintained
  - Secure
  - Community
- Make images as small as possible:
  - Start with minimal base images (ie: alpine linux)
  - Reduce the number of layers when possible (ie: combine RUN instructions)
- Make images build fast:
  - Each instruction creates an intermediate layer in your cache
  - Order your Dockerfile instructions and start with the layers (instructions) which rarely change.
- Images should be pre-configured but customizable (as much as possible)
  - Easy to start
  - Easy to override specific configurations
- Use the .dockerignore file to exclude irrelevant files

Docker Compose

1. Command Line Interface tool

2. Define and run multi Docker-container systems

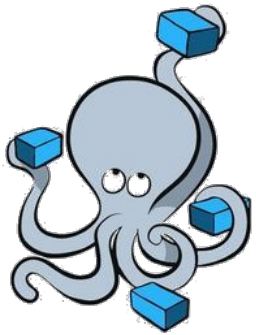

3. No more tedious docker run commands

4. Easy to run reproductible / complex systems

Docker Compose

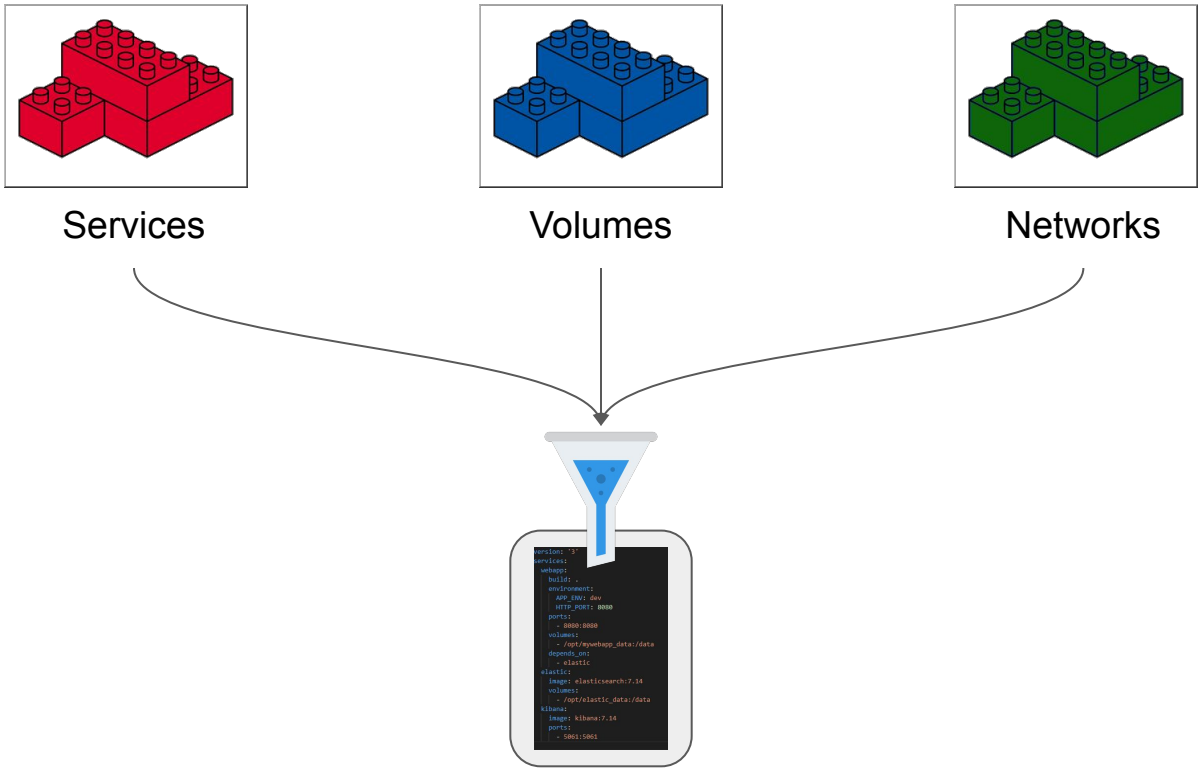

Docker Compose

- IaC : Describe resources & services and how they integrate with each other in a single Yaml file (containers, volumes, networks).

- Very useful for development environments to start / share / operate multi-container applications.

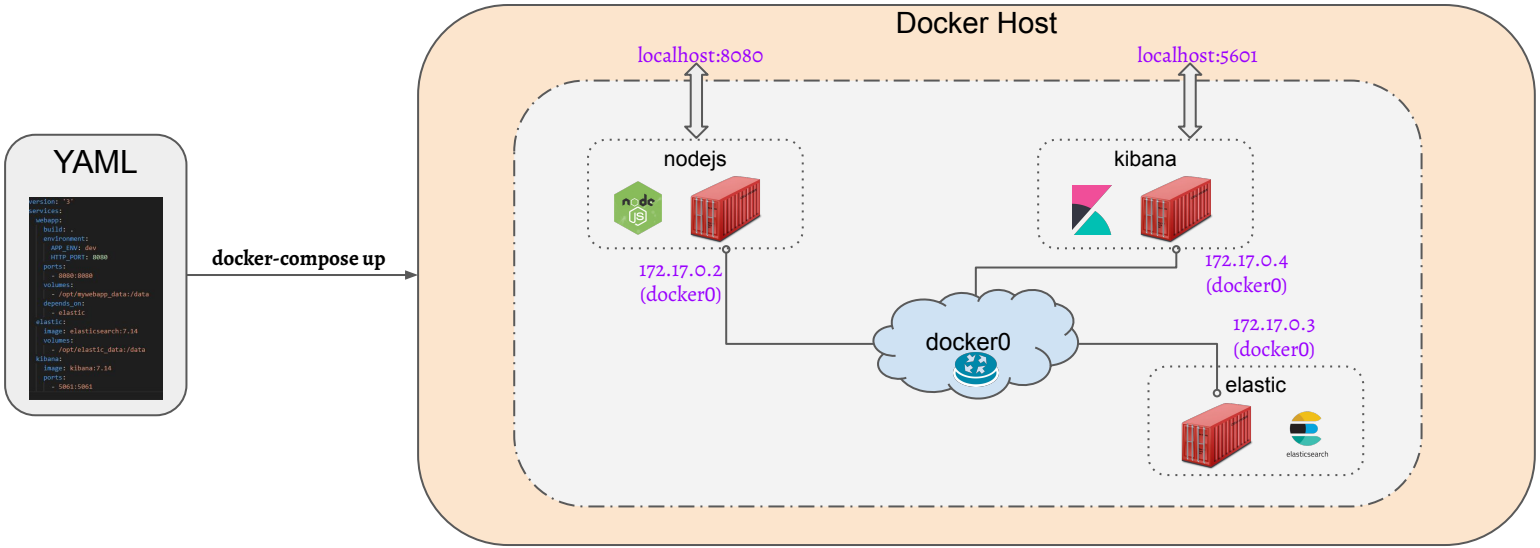

Docker Compose

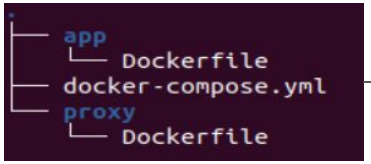


```yaml
version: "3.3"

services:
  proxy:
    build: ./proxy
    ports:
      - 8081:80
    networks:
      - frontend
  app:
    build: ./app
    environment:
      DB_USER: postgres
      DB_PASSWORD: admin
    networks:
      - frontend
      - backend
  db:
    image: postgres
    environment:
      POSTGRES_USER: postgres
      POSTGRES_PASSWORD: admin
    volumes:
      - db-vol
    networks:
      - backend

volumes:
  db-vol:

networks:
  frontend:
  backend:
```

Docker Compose

- Interesting to use variables in compose files:
  - Instead of repeating values across the files (ie: DB user/password)
  - To externalize values which may change (ie: image versions)

- Compose files can include variables. This allows us to inject configuration when starting services.

- This can be achieved with .env files

```
version: '3'
services:
  db:
    image: mysql:latest
    environment:
      - MYSQL_DATABASE=${MYSQL_DB}
      - MYSQL_USER=${MYSQL_USER}
      - MYSQL_PASSWORD=${MYSQL_PW}
```

docker-compose.yml

```
MYSQL_DB=db
MYSQL_USER=my_username
MYSQL_PW=K6c574T@9Qd
```

.env

```
version: '3'
services:
  db:
    image: mysql:latest
    env_file: .env
```

docker-compose.yml

```
MYSQL_DATABASE=db
MYSQL_USER=my_username
MYSQL_PASSWORD=K6c574T@9Qd
```

.env

Docker Swarm

- Docker Compose : Tool used to manage multi-container applications on a **single machine** only.

- Docker Swarm : Tool designed to manage multi-container applications on **clusters**. Acts as a container **orchestrator**.

- Both Compose & Swarm can use a same docker-compose file.

- Some features may be limited to either Compose or Swarm.
  For instance, secrets may be defined in Swarm but would be ignored in Compose.

Docker Swarm